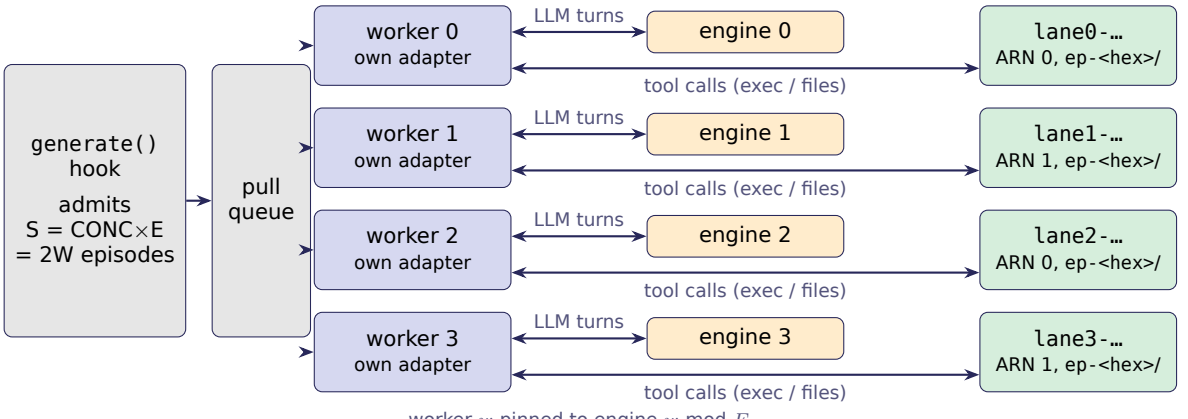
slime
GRPO trainer
Harness tier
W warm workers
Inference tier
SGLang, TP1, GPUs 4–7
Sandbox tier
AgentCore microVMs
generate() hook
admits S = CONC×E = 2W episodes
pull queue
worker 0
own adapter
worker 1
own adapter
worker 2
own adapter
worker 3
own adapter
LLM turns
engine 0
engine 1
engine 2
engine 3
tool calls (exec / files)
lane0-…
ARN 0, ep-<hex>/
lane1-…
ARN 1, ep-<hex>/
lane2-…
ARN 0, ep-<hex>/
lane3-…
ARN 1, ep-<hex>/
worker w pinned to engine w mod E
(bypasses the router's per-request LB)
opened once per worker, kept for the run;
fresh workdir per episode, rm -rf after