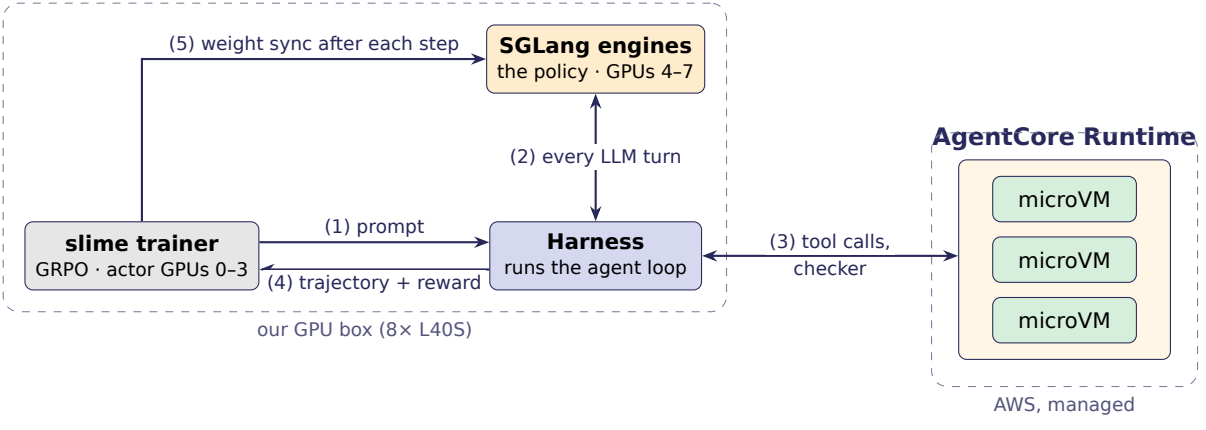
(5) weight sync after each step
SGLang engines
the policy · GPUs 4–7
(2) every LLM turn
slime trainer
GRPO · actor GPUs 0–3
(1) prompt
(4) trajectory + reward
Harness
runs the agent loop
(3) tool calls,
checker
AgentCore Runtime
microVM
microVM
microVM
our GPU box (8× L40S)
AWS, managed