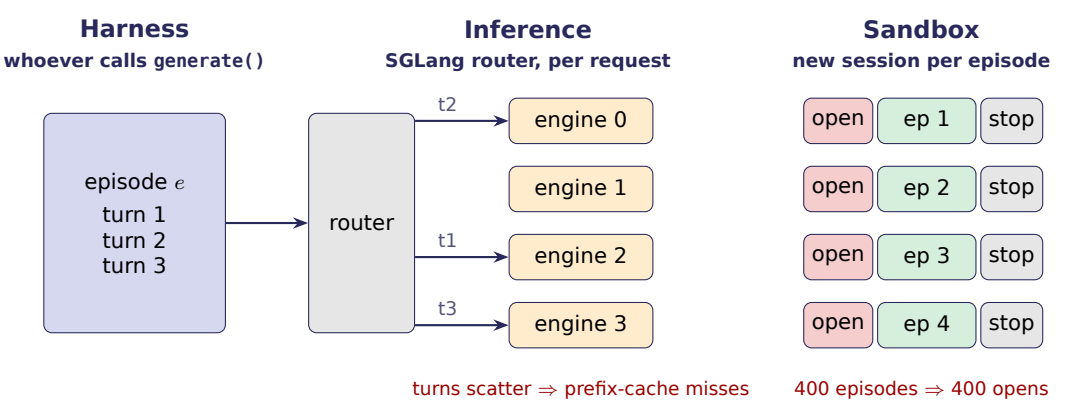
Harness
whoever calls generate()
Inference
SGLang router, per request
Sandbox
new session per episode
episode e
turn 1
turn 2
turn 3
router
t2
t1
t3
engine 0
engine 1
engine 2
engine 3
open
ep 1
stop
open
ep 2
stop
open
ep 3
stop
open
ep 4
stop
turns scatter ⇒ prefix-cache misses
400 episodes ⇒ 400 opens