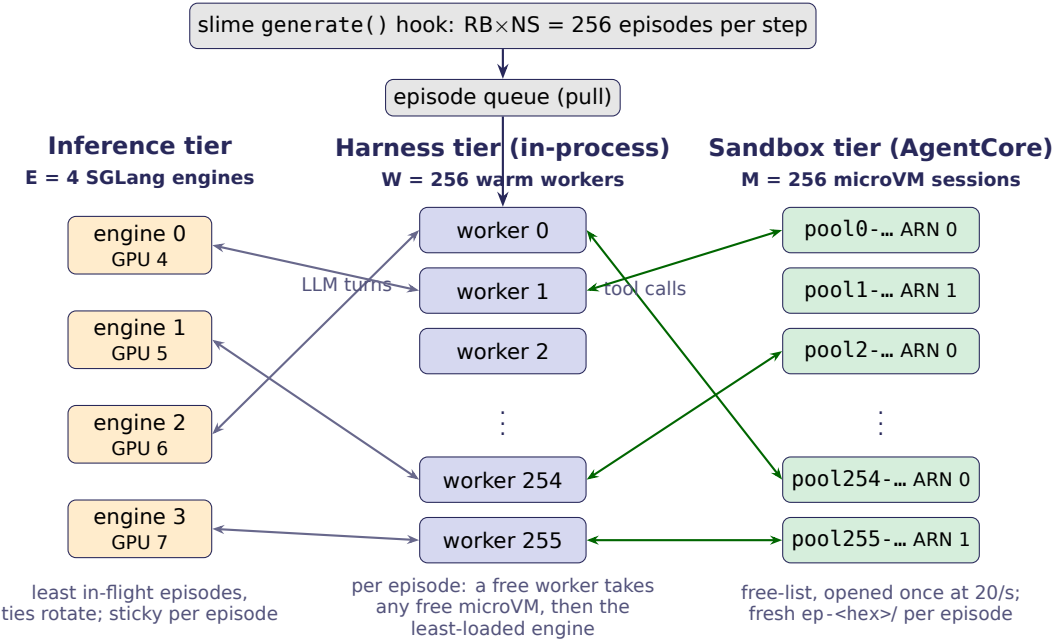
slime generate() hook: RB×NS = 256 episodes per step
episode queue (pull)
Inference tier
E = 4 SGLang engines
Harness tier (in-process)
W = 256 warm workers
Sandbox tier (AgentCore)
M = 256 microVM sessions
engine 0
GPU 4
engine 1
GPU 5
engine 2
GPU 6
engine 3
GPU 7
LLM turns
worker 0
worker 1
worker 2
worker 254
worker 255
tool calls
pool0-… ARN 0
pool1-… ARN 1
pool2-… ARN 0
pool254-… ARN 0
pool255-… ARN 1
least in-flight episodes,
ties rotate; sticky per episode
per episode: a free worker takes
any free microVM, then the
least-loaded engine
free-list, opened once at 20/s;
fresh ep-<hex>/ per episode