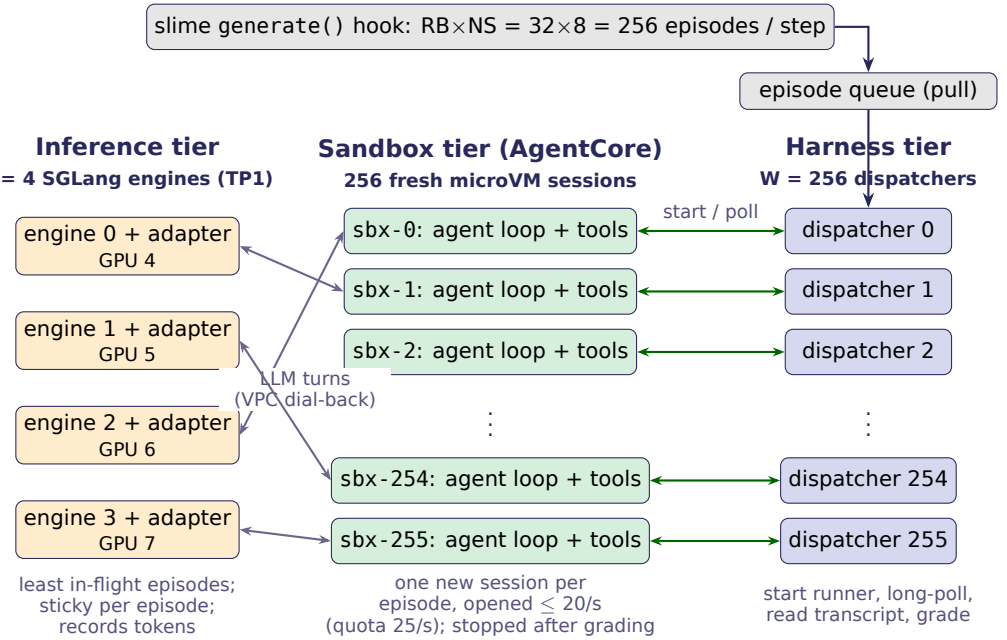
slime generate() hook: RB×NS = 32×8 = 256 episodes / step
episode queue (pull)
Inference tier
E = 4 SGLang engines (TP1)
Sandbox tier (AgentCore)
256 fresh microVM sessions
Harness tier
W = 256 dispatchers
start / poll
engine 0 + adapter
GPU 4
engine 1 + adapter
GPU 5
engine 2 + adapter
GPU 6
engine 3 + adapter
GPU 7
LLM turns
(VPC dial-back)
sbx-0: agent loop + tools
sbx-1: agent loop + tools
sbx-2: agent loop + tools
sbx-254: agent loop + tools
sbx-255: agent loop + tools
dispatcher 0
dispatcher 1
dispatcher 2
dispatcher 254
dispatcher 255
least in-flight episodes;
sticky per episode;
records tokens
one new session per
episode, opened ≤ 20/s
(quota 25/s); stopped after grading
start runner, long-poll,
read transcript, grade